



A typed-in wrong answer can keep driving the model — even when it looks like it revised

Takeaway: that stored answer state (hidden activations at answer tokens, middle layers) causally controls W–C preference — including on Qwen, which already looks revised.


Panel A — Claim: a typed wrong answer is stored inside the model and reused later




Panel B — Causal test: replace that stored state, and the model's W-vs-C preference shifts (± 4.9)^[3]



Panel C1 — Unmodified, GPT-2 sticks to W; Qwen chooses C



Panel C2 — Overwriting that stored state still steers Qwen (-2.1); random tokens don't



Notes

[1] W / C — wrong vs correct answer token in the prompt. First marked in Panel A.

[2] Stored answer state — the model's hidden activations at the typed-answer token positions, in middle layers (not input/output; GPT-2 $\approx$ L5–8, Qwen $\approx$ L8–23). First marked in Panel A; Panel B edits this box. Panel C2 "rand. pos. (localization)" — same overwrite at unrelated tokens; null effect shows control is site-specific.

[3] Δ — change in $\text{logit}(W) - \text{logit}(C)$ after the edit (after – before). Negative = shifted toward C. Cited on Panel B's title; Panel C2 uses the same Δ .

[4] Score — current $\text{logit}(W) - \text{logit}(C)$, not a change. >0 prefers W; <0 prefers C.